

Linear Algebra

Linear Dependence

Consider a set of k n -dimensional vectors $\{X_1, X_2, \dots, X_k\}$. These vectors are,

Definition 0.1. *linearly dependent* if there exists a set of scalars $\{a_1, a_2, \dots, a_k\}$ such that

$$a_1X_1 + a_2X_2 + \dots + a_kX_k = 0$$

where at least one $a_i \neq 0$.

Alternatively, they are,

Definition 0.2. *linearly independent* if the only set of scalars $\{a_1, a_2, \dots, a_k\}$ that satisfies the above condition is $a_1, a_2, \dots, a_k = 0$.

If we collect these k column-vectors in a matrix, $X = [X_1 \ X_2 \ \dots \ X_k]$, then the linear dependence condition can be written as,

$$a_1X_1 + a_2X_2 + \dots + a_kX_k = [X_1 \ X_2 \ \dots \ X_k] \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_k \end{bmatrix} = Xa = 0$$

Given any $n \times k$ matrix X , its columns are,

Definition 0.3. *linearly dependent* if there exists a vector $a \in \mathbb{R}^k$ such that $a \neq 0$ and $Xa = 0$;

or,

Definition 0.4. *linearly independent* if the only vector $a \in \mathbb{R}^k$ such that $Xa = 0$ is $a = 0$.

For any matrix there may be more than one vector $a \in \mathbb{R}^k$ such that $Xa = 0$. Indeed, if both $a_1, a_2 \in \mathbb{R}^k$

satisfy this condition and $a_1 \neq a_2$ then you can show that any linear combination of $\{a_1, a_2\}$ satisfies the

condition $X(a_1b_1 + a_2b_2) = 0$ for $b_1, b_2 \in \mathbb{R}$. Thus, there exists an entire set of vectors which satisfy this condition. This set is referred to as the,

Definition 0.5. *null space* of X ,

$$\mathcal{N}(X) = \{a \in \mathbb{R}^k : Xa = 0\}$$

It should be evident from the definition that if the columns of X are linearly independent then $\mathcal{N}(X) = \{0\}$, a singleton. That is, it just includes the 0-vector.

Vector spaces, bases, and spans

Here, we concern ourselves only with real vectors from \mathbb{R}^n .

Definition 0.6. A *vector space*, denoted \mathcal{V} , refers to a set of vectors which is closed under finite addition and scalar multiplication.

Definition 0.7. A set of k linearly independent vectors, $\{X_1, X_2, \dots, X_k\}$, forms a *basis* for vector space \mathcal{V} if $\forall y \in \mathcal{V}$ there exists a set of k scalars such that,

$$y = X_1 b_1 + X_2 b_2 + \dots + X_k b_k$$

Based on these definitions, it is evident that for the Euclidean space, \mathbb{E}^n , any n linearly independent vectors from \mathbb{R}^n is a basis. For example, any point in \mathbb{E}^2 can be defined as a multiple of,

$$\begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Consider again the $n \times k$ matrix X , where $k < n$. Then we define the,

Definition 0.8. *column space* (or *span*) of X , denoted $\mathcal{S}(X)$, as the vector space generate by the k columns of X . Formally,

$$\mathcal{S}(X) = \{y \in \mathbb{R}^n : y = Xb \quad \text{for some } b \in \mathbb{R}^k\}$$

A property to note about the span or column space X is,

Result: $\mathcal{S}(X) = \mathcal{S}(XX') :::$

where XX' is a $n \times n$ matrix.

Finally, we can define the,

Definition 0.9. *orthogonal column space* (or *orthogonal span*) of X as,

$$\mathcal{S}^\perp(X) = \{y \in \mathbb{R}^k : y'x = 0 \quad \forall x \in \mathcal{S}(X)\}$$

Rank

Consider a $n \times k$ matrix X , the

Definition 0.10. *row rank* of X is the maximum number of linearly independent rows:

$$\text{rowrank}(X) \leq n$$

We say that matrix X has *full* row rank if $\text{rowrank}(X) = n$.

The,

Definition 0.11. *column rank* of X is the maximum number of linearly independent columns:

$$\text{colrank}(X) \leq k$$

We say that matrix X has *full* column rank if $\text{colrank}(X) = k$.

An important result is,

- **Result:** the rank of X :

$$r(X) = \text{rowrank}(X) = \text{colrank}(X) \Rightarrow r(X) \leq \min\{n, k\}$$

In addition, since the $r(X)$ depends on the number of linearly independent columns, we can say that,

- **Result:** the dimension of $\mathcal{S}(X)$, $\dim(\mathcal{S}(X))$, is given by the $r(X)$.

Here are a few additional results,

- **Result:** $r(X) = r(X')$
- **Result:** $r(XY) \leq \min\{r(X), r(Y)\}$
- **Result:** $r(XY) = r(X)$ if Y is square and full rank
- **Result:** $r(X + Y) \leq r(X) + r(Y)$

Properties of square matrices

Consider the case of a square, $n \times n$, matrix A . We say that,

Definition 0.12. A is *singular* if the $r(A) < n$,

or that,

Definition 0.13. A is *non-singular* if the $r(A) = n$.

The singularity of a square matrix is important as it determines the invertibility of a matrix, which typically relates the existence of a unique solution in systems of linear equations. Here are a few key results,

- **Result:** There exists a matrix $B = A^{-1}$, such that $AB = I_n$ (where I_n is the identity matrix), if and only if A is non-singular.
- **Result:** A is non-singular if and only if the determinant of A is non-zero: $\det(A) \neq 0$.¹
- **Result:** Likewise, A is singular if and only if $\det(A) = 0$.
- **Result:** $AA^{-1} = A^{-1}A = I$
- **Result:** $(A')^{-1} = (A^{-1})'$
- **Result:** If their respective inverses exist, then $(AB)^{-1} = B^{-1}A^{-1}$.
- **Result:** $\det(AB) = \det(A)\det(B)$
- **Result:** $\det(A^{-1}) = \det(A)^{-1}$

For any square matrix A ,

Definition 0.14. the *trace* of A is the sum of all diagonal elements:

$$tr(A) = \sum_{i=1}^n a_{ii}$$

Regarding the trace of a square matrix, here are a few important results:

- **Result:** $tr(A + B) = tr(A) + tr(B)$
- **Result:** $tr(\lambda A) = \lambda tr(A)$ where λ is a scalar
- **Result:** $tr(A) = tr(A')$

¹These notes do not cover how to calculate the determinant of a square matrix. You should be able to find a definition easily online.

- **Result:** $\text{tr}(AB) = \text{tr}(BA)$ where AB and BA are both square, but need not be of the same order.
- **Result:** $\|A\| = (\text{tr}(A'A))^{1/2}$

Properties of symmetric matrices

A symmetric matrix has the property that $A = A'$. Therefore, A must be square.

Here are a few important results concerning symmetric matrices.

- **Result:** A^{-1} exists if $\det(A) \neq 0$ and $r(A) = n$
- **Result:** A is *diagonalizable*.²
- **Result:** The eigenvector decomposition of a square matrix gives you $A = C\Lambda C^{-1}$ where Λ is a diagonal matrix of eigenvalues and C a matrix of the corresponding eigenvectors. The symmetry of A gives you that $C^{-1} = C' \Rightarrow A = C\Lambda C'$ with $C'C = CC' = I_n$.³

A key definition concerning symmetric matrices is their positive definiteness:

Definition 0.15. A is *positive semi-definite* if for any $x \in \mathbb{R}^n$, $x'Ax \geq 0$.

Given the eigenvector decomposition of a symmetric matrix, *positive semi-definiteness* implies Λ is *positive semi-definite*: $\lambda_i \geq 0 \quad \forall i$. Likewise,

Definition 0.16. A is *positive definite* if for any $x \in \mathbb{R}^n$, $x'Ax > 0$.

Again, based on the eigenvector decomposition, *positive semi-definiteness* implies Λ is *positive definite*: $\lambda_i > 0 \quad \forall i$.

A few more results are:

- **Result:** $\text{tr}(A) = \sum_{i=1}^n \lambda_i$
- **Result:** $r(A) = r(\Lambda)$
- **Result:** $\det(A) = \prod_{i=1}^n \lambda_i$

²A matrix is diagonalizable if it is *similar* to some other diagonal matrix. Matrices B and C are similar if $C = PBP^{-1}$. A square matrix which is not diagonalizable is *defective*. This property relates closely to eigenvector decomposition.

³Recall, an eigenvalue and eigenvector pair, (λ, c) , of matrix A satisfy:

$$Ac = \lambda c \Rightarrow (A - \lambda I_n)c = 0$$

This last result can be used to prove that any positive definite matrix is non-singular and therefore has an inverse.

Any full-rank, positive semi-definite, symmetric matrix B has the additional properties:

- **Result:** $B = C\Lambda C'$ and $B^{-1} = C\Lambda^{-1}C'$
- **Result:** We can define the square-root of B as $B^{1/2} = C\Lambda^{1/2}C'$. Similarly, $B^{-1/2} = C\Lambda^{-1/2}C'$.

Properties of idempotent matrices

An idempotent matrix has the property that $D = DD$. Therefore, D must be square.

Here are a few important results concerning idempotent matrices.

- **Result:** D is positive definite
- **Result:** D is diagonalizable
- **Result:** $(I_n - D)$ is also an idempotent matrix
- **Result:** With the exception of I_n , all idempotent matrices are singular.
- **Result:** $r(D) = tr(D) = \sum_{i=1}^n \lambda_i$
- **Result:** $\lambda_i \in \{0, 1\} \quad \forall i$

Projection matrices are idempotent, but need not be symmetric. However, for the purposes of this module we will deal exclusively with symmetric idempotent projection matrices.

Vector Differentiation

Here we will look at the derivatives of scalar with respect to (W.r.t.) a vector. You can also define other derivatives, such as the derivative of a vector w.r.t. a vector and the derivative of a scalar with respect to a matrix. However, these are not needed for these notes.

General case

Suppose $f(x) \in R$ (i.e. a scalar) and $x \in R^n$ (i.e. a $n \times 1$ vector). Then we can define the partial derivative of $f(x)$ w.r.t. x as,

$$\frac{\partial f(x)}{\partial x} = \begin{bmatrix} \frac{\partial f(x)}{\partial x_1} \\ \frac{\partial f(x)}{\partial x_2} \\ \vdots \\ \frac{\partial f(x)}{\partial x_n} \end{bmatrix}$$

Linear: scalar case

A special case is when $f(x)$ is linear in x ,

$$f(x) = a'x = \sum_{i=1}^n a_i x_i$$

for $a \in R^n$. The derivative of $a'x$ with respect to the **vector** x can be defined as,

$$\begin{aligned} \frac{\partial a'x}{\partial x} &= \begin{bmatrix} \frac{\partial a'x}{\partial x_1} \\ \frac{\partial a'x}{\partial x_2} \\ \vdots \\ \frac{\partial a'x}{\partial x_n} \end{bmatrix} \\ &= \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix} \\ &= a \end{aligned}$$

since the the partial derivate of $a'x = \sum_{i=1}^n a_i x_i$ w.r.t. x_i is just the scalar a_i .

The derivative of a scalar w.r.t. to a vector yields a vector of partial derivatives.

Since $a'x$ is a scalar, it is by definition symmetric: $a'x = x'a$. Thus,

$$\frac{\partial x'a}{\partial x} = \frac{\partial a'x}{\partial x} = a$$

Linear: vector case

Suppose $f(x)$ is a linear transformation of x ,

$$f(x) = A'x$$

for any $m \times n$ matrix A ,

$$A = \begin{bmatrix} a'_1 \\ a'_2 \\ \vdots \\ a'_m \end{bmatrix}$$

where $a_i \in R^n \forall i = 1, \dots, m$ and,

$$Ax = \begin{bmatrix} a'_1 x \\ a'_2 x \\ \vdots \\ a'_m x \end{bmatrix}$$

Note, $f(x) = Ax \in R^m$, a $m \times 1$ vector. We can then define,

$$\begin{aligned} \frac{\partial Ax}{\partial x'} &= \begin{bmatrix} \frac{\partial a'_1 x}{\partial x_1} & \frac{\partial a'_1 x}{\partial x_2} & \cdots & \frac{\partial a'_1 x}{\partial x_n} \\ \frac{\partial a'_2 x}{\partial x_1} & \frac{\partial a'_2 x}{\partial x_2} & \cdots & \frac{\partial a'_2 x}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial a'_m x}{\partial x_1} & \frac{\partial a'_m x}{\partial x_2} & \cdots & \frac{\partial a'_m x}{\partial x_n} \end{bmatrix} \\ &= \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix} \\ &= A \end{aligned}$$

Since Ax is $m \times 1$ column vector, we take the derivative w.r.t. x' a row vector and not the column vector x . This results in a matrix of partial derivatives.

Quadratic form

A second special case is where the function takes on the quadratic form,

$$f(x) = x'Ax = \sum_{i=1}^N \sum_{j=1}^n a_{ij}x_i x_j$$

for $n \times n$ (square) matrix A . As in the first linear case, $f(x)$ is scalar.

Define $c = Ax$, the $x'Ax = x'c$. From the linear case, we know that,

$$\frac{\partial x'c}{\partial x} = c$$

Similarly, if we define $d = A'x$ then $x'Ax = d'x$. From the linear case, we know that,

$$\frac{\partial d'x}{\partial x} = d$$

We can define the total derivative as the sum of the partial derivatives w.r.t. to the first and second x . Combining these two results, we have that,

$$\frac{\partial x'Ax}{\partial x} = Ax + A'x$$

And **if** A is symmetric, this result simplifies to $2Ax$.